



Joint Tech Internship Community Program

Assignment 1

SUBMITTED BY

HEMANATH B

(hemanathb10@gmail.com)

Sample Dataset : House Price Prediction

House ID	Square Feet	No of Bedrooms	No of Bathrooms	Age of House (years)	Price (\$)
1	1500	3	2	10	300,000
2	2000	4	3	5	450,000
3	1200	2	1	15	200,000
4	1800	3	2	8	350,000
5	2500	4	3	3	500,000
6	1000	2	1	20	150,000
7	1600	3	2	12	320,000
8	2200	4	3	6	480,000

This dataset can be used to illustrate the following terminologies:

1) Feature:

- **Definition:** Features are individual measurable properties or characteristics of a phenomenon being observed.
- **Example in Dataset:** Square Feet, Number of Bedrooms, Number of Bathrooms, Age of House are features of the houses.

2) Label:

- **Definition:** The label is the output variable that we are trying to predict or classify.
- **Example in Dataset:** The Price (\$) is the label we want to predict.

3) Prediction:

- **Definition:** A prediction is the output of a machine learning model when it is given an input example.
- **Example in Dataset:** Predicting the price of a house based on its features.

4) Outlier:

- **Definition:** An outlier is an observation point that is distant from other observations.
- **Example in Dataset:** If most houses are priced around \$300,000 to \$500,000, a house priced at \$150,000 might be considered an outlier.

5) Test Data:

- **Definition:** Test data is a subset of the dataset used to evaluate the performance of a trained model.
- **Example in Dataset:** 2 - 3 records from the dataset can be set aside as test data to evaluate the model's performance after training.

6) Training Data:

- **Definition:** Training data is a subset of the dataset used to train the model.
- **Example in Dataset:** 4 - 5 records in the dataset would be used as training data to teach the model the relationship between the features and the label.

7) Model:

- **Definition:** A model is a mathematical representation of a real-world process. In machine learning, it is trained to make predictions.
- **Example in Dataset:** A regression model that predicts house prices based on their features.

8) Validation Data:

- **Definition:** Validation data is a subset of the dataset used to provide an unbiased evaluation of a model fit on the training dataset while tuning model hyperparameters.
- **Example in Dataset:** 1 – 2 records from the training data is used to validate the model's performance.

9) Hyperparameter:

- **Definition:** Hyperparameters are configuration settings used to tune how the machine learning model is trained.
- **Example in Dataset:** Examples include the learning rate, number of epochs, or the regularization parameter.

10) Epoch:

- **Definition:** An epoch is one complete pass through the entire training dataset.
- **Example in Dataset:** If the dataset is passed through the model 100 times during training, that would be 100 epochs.

11) Loss Function:

- **Definition:** A loss function measures how well the model's predictions match the true labels.
- **Example in Dataset:** Mean Squared Error (MSE) could be used as a loss function to measure the difference between predicted house prices and actual prices.

12) Learning Rate:

- **Definition:** The learning rate is a hyperparameter that controls how much to change the model in response to the estimated error each time the model weights are updated.
- **Example in Dataset:** Setting a learning rate to 0.01 for updating the model weights during training.

13) Overfitting:

- **Definition:** Overfitting occurs when a model learns the training data too well, including noise and details, leading to poor performance on new data.
- **Example in Dataset:** If the model performs exceptionally well on the training houses but poorly on the test houses, it may be overfitting.

14) Underfitting:

- **Definition:** Underfitting occurs when a model is too simple to capture the underlying pattern of the data.
- **Example in Dataset:** If the model performs poorly on both training and test data, it may be underfitting.

15) Regularization:

- **Definition:** Regularization techniques are used to reduce the risk of overfitting by adding a penalty to the loss function for large coefficients.
- **Example in Dataset:** Applying L2 regularization to penalize large weights in the model.

16) Cross-Validation:

- **Definition:** Cross-validation is a technique for assessing how the results of a statistical analysis will generalize to an independent dataset.
- **Example in Dataset:** Using k-fold cross-validation to divide the dataset into k parts and training the model k times, each time using a different part as the validation set.

17) Feature Engineering:

- **Definition:** Feature engineering involves creating new features or modifying existing ones to improve model performance.
- **Example in Dataset:** Creating a new feature such as Price per Square Foot by dividing the price by the square feet.

18) Dimensionality Reduction:

- **Definition:** Dimensionality reduction is the process of reducing the number of random variables under consideration.
- **Example in Dataset:** Using Principal Component Analysis (PCA) to reduce the number of features while retaining most of the informations in the data.

19) Bias:

- **Definition:** Bias is an error due to overly simplistic assumptions in the learning algorithm.
- **Example in Dataset:** If the model consistently predicts house prices lower than the actual prices, it may have a bias.

20) Variance:

- **Definition:** Variance is an error due to too much complexity in the learning algorithm.
- **Example in Dataset:** If the model's predictions vary widely for similar houses in the test set, it may have high variance.